

## Week 1: Introduction to Generative AI and Large Language Models (LLMs)

### Overview

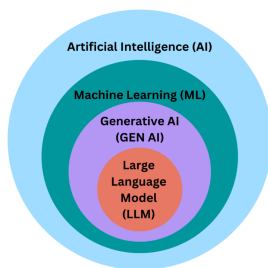
This week, I have learned **Generative AI** and **Large Language Models (LLMs)**. Moreover, I have also focused on how LLMs work, their key use cases, and the challenges involved in training them. The topics covered include:

- **Transformer architecture** – the foundation of modern LLMs.
  - **Text generation techniques** – how LLMs create human-like text.
  - **Memory optimization** – ways to manage computational resources efficiently.
  - **Scaling models across GPUs** – how to train and deploy large models.
- 

### What is Generative AI?

Generative AI refers to models that create new content, such as text, images, or code, by mimicking human output.

**Large Language Models (LLMs)** are a subset of generative AI specifically designed for working with text. These models are trained on vast amounts of text data to recognize patterns and generate contextually relevant responses.



### Example of How LLMs Work:

When given a sentence starter like **"The sky is..."**, an LLM predicts the next word based on its training data, generating something like **"blue."**

---

### Common Use Cases of LLMs

LLMs can perform various text-based tasks, including:

#### 1. Text Generation

LLMs can write entire articles, essays, or even code.

- **Example:** A chatbot responding with, "How can I assist you today?"

#### 2. Translation

They convert text between languages accurately.

- **Example:** "Hello" → "Hola" (Spanish)

#### 3. Summarization

LLMs can condense long articles into short summaries.

- **Example:** Summarizing a 1,000-word news article into three sentences.

#### 4. Named Entity Recognition (NER)

This identifies important entities (names, dates, places, etc.).

- **Example:** "Steve works at Google" → **Person:** Steve, **Organization:** Google.

---

## Text Generation Before Transformers

Before **Transformers**, earlier models used **Recurrent Neural Networks (RNNs)**, which processed text sequentially (word-by-word). However, RNNs had limitations:

### 1. Short-Term Memory Issues:

- RNNs struggled with long sentences.
- Example: "The cat, which was chased by the dog, ran into the garden." By the time it reached "garden," the model might forget the "cat."

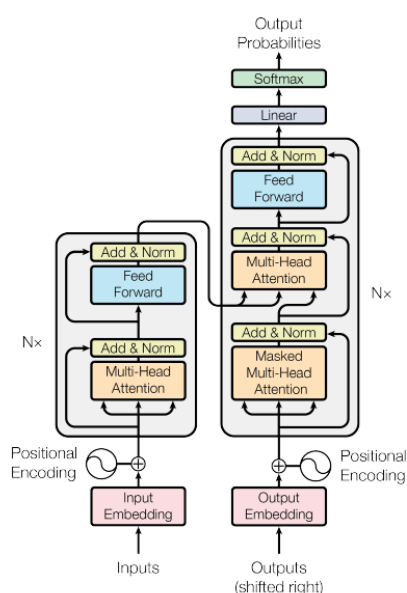
### 2. Slow Training Speed:

- RNNs processed one word at a time, making them inefficient.

---

## The Transformer Architecture

Introduced in **2017**, **Transformers** revolutionized text generation by allowing models to process all words simultaneously instead of sequentially.



### Key Components of a Transformer:

#### 1. Tokenization: Converts words into numerical tokens.

- Example: "Hello" → [23, 45]

#### 2. Embeddings: Maps tokens to numerical vectors to represent meaning.

- Example: "King" → [0.3, -0.2, 0.5]

#### 3. Positional Encoding: Ensures word order is considered.

- Example: "Dog bites man" vs. "Man bites dog" (different meanings)

#### 4. Self-Attention Mechanism: Helps understand

relationships between words.

- Example: "The teacher gave the students homework" → "teacher" relates to "students" and "homework."

#### 5. Feed-Forward Network: Converts self-attention outputs into probabilities for predicting the next word.

---

## How Transformers Generate Text

When you ask a model to generate a story, the process involves:

1. **Input:** "Write a story about a robot."
2. **Tokenization:** Converts input into numbers like [10, 23, 45].
3. **Processing:** Uses self-attention to predict the next words.
4. **Output:** "Once, a robot named Zeta explored Mars and discovered ancient ruins."

### Controlling Output with Parameters

- **Temperature:** Adjusts randomness.
    - **Low (0.1):** Predictable, structured output. Example: "The cat sat on the mat."
    - **High (2.0):** More creative output. Example: "The cat tap-danced on the moon."
  - **Max Tokens:** Limits response length. Example: Setting **Max Tokens = 50** ensures the response doesn't exceed ~50 words.
- 

### Prompting and Prompt Engineering

Prompting techniques help guide LLM responses effectively:

- **Zero-Shot Prompting:** No examples provided.
    - Example: "Summarize this article: [text]."
  - **One-Shot Prompting:** One example provided.
    - Example: "Translate 'Hello' to French: Bonjour. Translate 'Goodbye' to French: \_\_\_\_"
  - **Few-Shot Prompting:** Multiple examples provided.
    - Example: "Review: 'Great movie!' → Positive. Review: 'Bad acting' → Negative. Review: [Your text] → \_\_\_\_"
- 

### Building a Generative AI Model: Project Lifecycle

1. **Define the Use Case:** Example: "Summarize legal documents."
  2. **Choose a Model:** Use an existing model (e.g., T5) or train a custom one.
  3. **Adapt the Model:** Fine-tune it with domain-specific data.
  4. **Deploy the Model:** Integrate into applications via APIs.
- 

### Lab 1: Dialogue Summarization with FLAN-T5

Goal: Summarize customer support conversations.

Steps:

Install necessary libraries:  
pip install transformers datasets

1. Load the model:  

```
from transformers import AutoTokenizer, AutoModelForSeq2SeqLM

tokenizer = AutoTokenizer.from_pretrained("google/flan-t5-base")

model = AutoModelForSeq2SeqLM.from_pretrained("google/flan-t5-base")
```
  2. Generate a summary:  

```
inputs = tokenizer("Summarize: " + conversation_text, return_tensors="pt")

outputs = model.generate(inputs["input_ids"], max_length=100)

summary = tokenizer.decode(outputs[0], skip_special_tokens=True)
```
- 

## Challenges in Training LLMs

1. **High GPU Memory Requirements:**
  - A 1B-parameter model requires **24GB** memory (FP32 precision).
2. **Optimization Strategies:**
  - **Quantization:** Reduces memory usage (FP32 → BFLOAT16 → INT8)
  - **Distributed Training:** Splits model across GPUs.
  - Pre-training large language models

Computational challenges of training LLMs  
Efficient multi-GPU compute strategies  
Scaling laws and compute-optimal models

**Pre-training for Domain Adaptation**

Pre-training from scratch is useful when a model needs to specialize in a particular field, such as finance, law, or medicine. In these cases, general-purpose models may not perform well due to a lack of domain-specific vocabulary or context understanding.

### When to Train a Model from Scratch

- **Specialized Vocabulary:** Fields like law and medicine use unique terms that general models might not understand well.
  - **Example:** A legal model needs to recognize terms like *res judicata* (a legal principle).
- **Highly Focused Data:** Some industries have proprietary or confidential data that general models lack access to.
  - **Example:** BloombergGPT was trained with 51% financial data to better understand financial reports and market trends.

### Benefits of Domain-Specific Pre-Training

- **Enhanced Understanding:** The model gains deeper insights into the domain-specific language.
- **Better Accuracy:** It provides more relevant and precise responses.

- **Improved Performance:** Even smaller models can outperform larger, general models if properly trained on high-quality, domain-specific data.
- 

## Computational Challenges of Training LLMs

### Key Issues in Training Large Models

1. **GPU Memory Constraints:**
  - Large models require vast amounts of memory.
  - Example: A model with **1 billion parameters** needs **24GB of memory** in FP32 precision.
2. **Optimizer States and Training Overhead:**
  - Advanced optimizers like Adam require **twice the model size** in memory.
  - Example: A **10B-parameter model** needs **80GB** for training (model + optimizer states).
3. **Quantization Techniques for Efficiency:**
  - **FP32 (32-bit floating point):** High precision but memory-intensive.
  - **BFLOAT16 (16-bit floating point):** Reduces memory usage while retaining numerical range.
  - **INT8 (8-bit integer):** Further reduces size but sacrifices some precision.

### Why Quantization Works

By using lower precision formats (BFLOAT16, INT8), models can significantly reduce memory usage without a major drop in performance. For example, **BFLOAT16 retains most of FP32's range** but uses **half the memory**, making it a preferred choice for training large-scale LLMs.

---

## Efficient Multi-GPU Compute Strategies

### 1. Distributed Data Parallel (DDP)

- Copies the entire model onto each GPU.
- Best for **small to medium models** (up to ~2B parameters).
- Example: Training **T5-3B** across 4 GPUs.

### 2. Fully Sharded Data Parallel (FSDP)

- Splits model parameters, optimizer states, and gradients across multiple GPUs.
- Uses **ZeRO Optimization** to reduce memory footprint.
- **ZeRO Stages:**
  - **Stage 1:** Shards optimizer states (4x memory savings).
  - **Stage 2:** Shards gradients (8x memory savings).
  - **Stage 3:** Shards model parameters (scales to 64+ GPUs).
- Example: Training **T5-11B** across **512 GPUs**.

### Trade-offs:

- **DDP** is simpler but uses more memory.
  - **FSDP** saves memory but increases GPU communication overhead.
- 

## Scaling Laws and Compute-Optimal Models

### Findings from the Chinchilla Paper

Research suggests that model performance **depends not just on size but on the amount of training data**.

- **Key Insight:** A model should be trained with **20 times more tokens than its number of parameters** for optimal performance.
  - **Example:** A **70B-parameter model** should be trained on **1.4 trillion tokens**.
- **Smaller, well-trained models can outperform larger ones** with the right amount of data.
  - **Example:** **Chinchilla-70B** outperforms **GPT-3-175B** because it was trained more efficiently.

### Practical Implications

- Instead of making models **bigger**, it is often better to train them **smarter** with **better data**.
- Many companies are shifting toward **high-quality, well-curated datasets** rather than just increasing model size.

Pre-training for domain adaptation this is missing

---

## Key Takeaways

- Transformers enable efficient and high-quality text generation.
- LLMs perform tasks like summarization, translation, and question-answering.
- Fine-tuning and prompting techniques improve model performance.
- Training LLMs require significant computational power.